



<https://youtu.be/9squ1uApOtY>

Gemini CLI Conceptual Architecture



Scott Rosenberg Fan Club

Callum McIntyre - **Group Leader** - Abstract, Architecture, Use Cases

Isaac Ouellette - **Presenter** - Use Cases, Slide Show, Script

Nicholas Saloufakos - **Presenter** - Use Cases, Derivation Process, Slideshow, Script

Jordan Bouckley - Architecture, AI Collaboration, Introduction, Overview

Andreea Cobzaru - Architecture, External Interfaces, Naming Conventions, Data Dictionary, Lessons Learnt, Conclusions

Zachary Gazaille - Architecture, AI Collaboration, References



What is Gemini CLI?

- Command line interface AI model
- Powered by Google Gemini LLM
- 3-tiered system | The CLI, The Core, The API
- Allows you to use AI within your command line/terminal
- Bridges the gap between your files and commands, and AI



Data Dictionary

- **Child Process:** an independent OS made by another process - executes shell commands in parallel
- **Context Window:** maximum amount of text the API can process
- **Event-Driven Model:** control flow pattern used by the system
- **Human-in-the-loop:** Safety design, human permission needed to execute tools
- **Non-deterministic:** property of LLMs where the same input can produce a different output
- **Promises:** unit of concurrency to handle asynchronous operations
- **Prompt engineering:** process performed by PromptConstructor, formats user input, history and instructions to optimize performance
- **ReAct pattern (Reason + Act):** logic flow used by the core, LLM makes a thought, uses a specific tool, observes the output and continues the reasoning
- **Selective context pruning:** performance optimization used by ContextManager, summarizes conversation history and removes irrelevant parts
- **Singleton pattern:** Architectural pattern to make sure a class has one instance
- **Streaming:** data transfer method, instead of waiting the API sends data in chunks
- **Token:** Smaller units that text is broken down into, so the LLM can process it
- **Tool:** function or script that the AI model can request to be executed.



Naming Conventions

- **CLI:** Command line interface - the user side
- **API:** Application protocol interface - the brain of the operation, ran within google on the server side
- **JSON:** JavaScript object notation
- **LLM:** Large language model
- **UI:** User interface



Architectural Styles

Style One: **Layered**

- 3 main layers | CLI - Core - API
- Organized hierarchically
- Each layer having specific set of info
 - Relies on other layers if not able to complete task
- Interacts only with layer adjacent
- Highly Modular → Evolvable

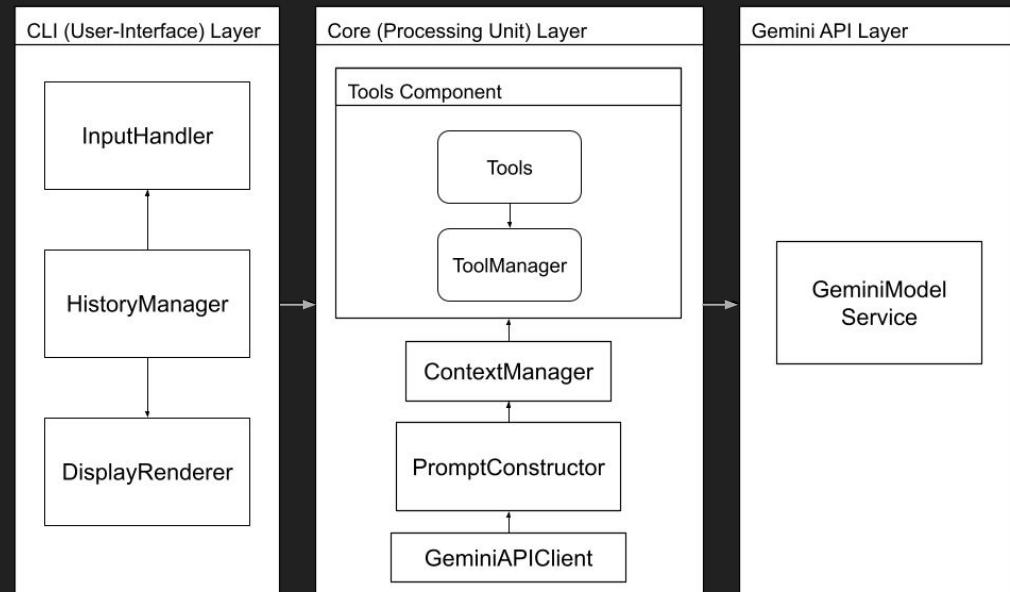
Style 2: **Client-Server**

- 2 Sides:
 - “Client” - CLI + Core
 - “Server” - API
- Client side takes and manages input
- Server side the brain of the operation
 - User prompts, CLI + Core contextualizes, sends to API
 - API sends structured response
- Logic held on server level, so local updates not usually needed
- Very low strain on local devices



Box-and-Arrow Diagram

- CLI → Core → API
- CLI handles the input, user history, and display
- Core houses the tools
- Creates context with user history
- Prompts generated with PromptConstructor
- GeminiAPIClient sends prompts to server
- GeminiModelService Does the deep thinking, reasoning, and decides what tools to use



Legend

→ Depends on



Component

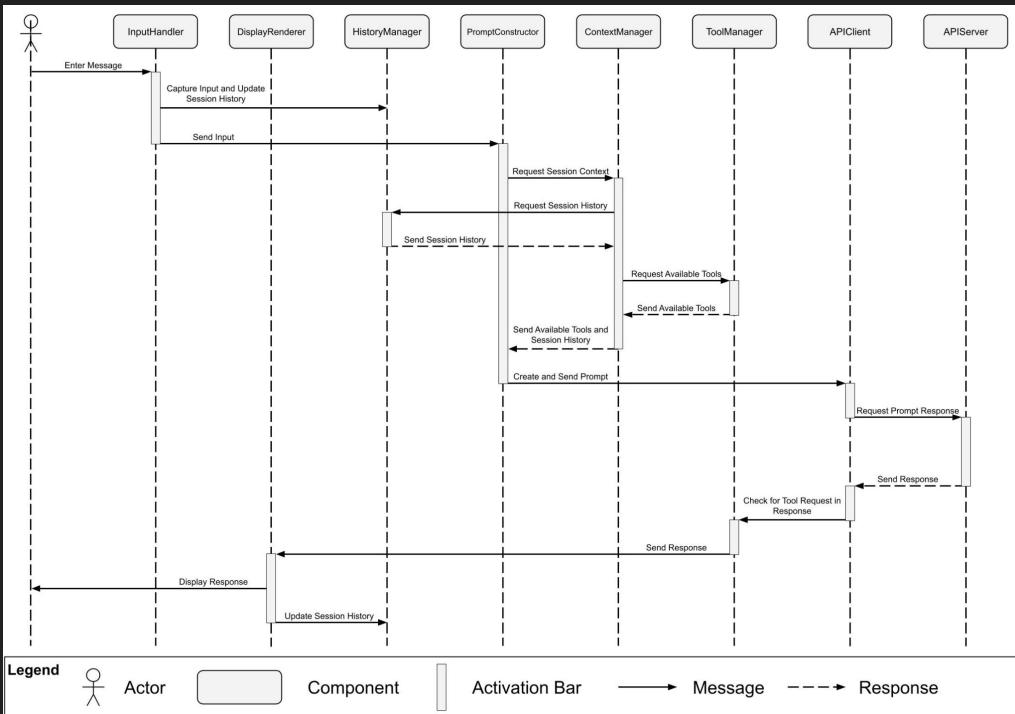


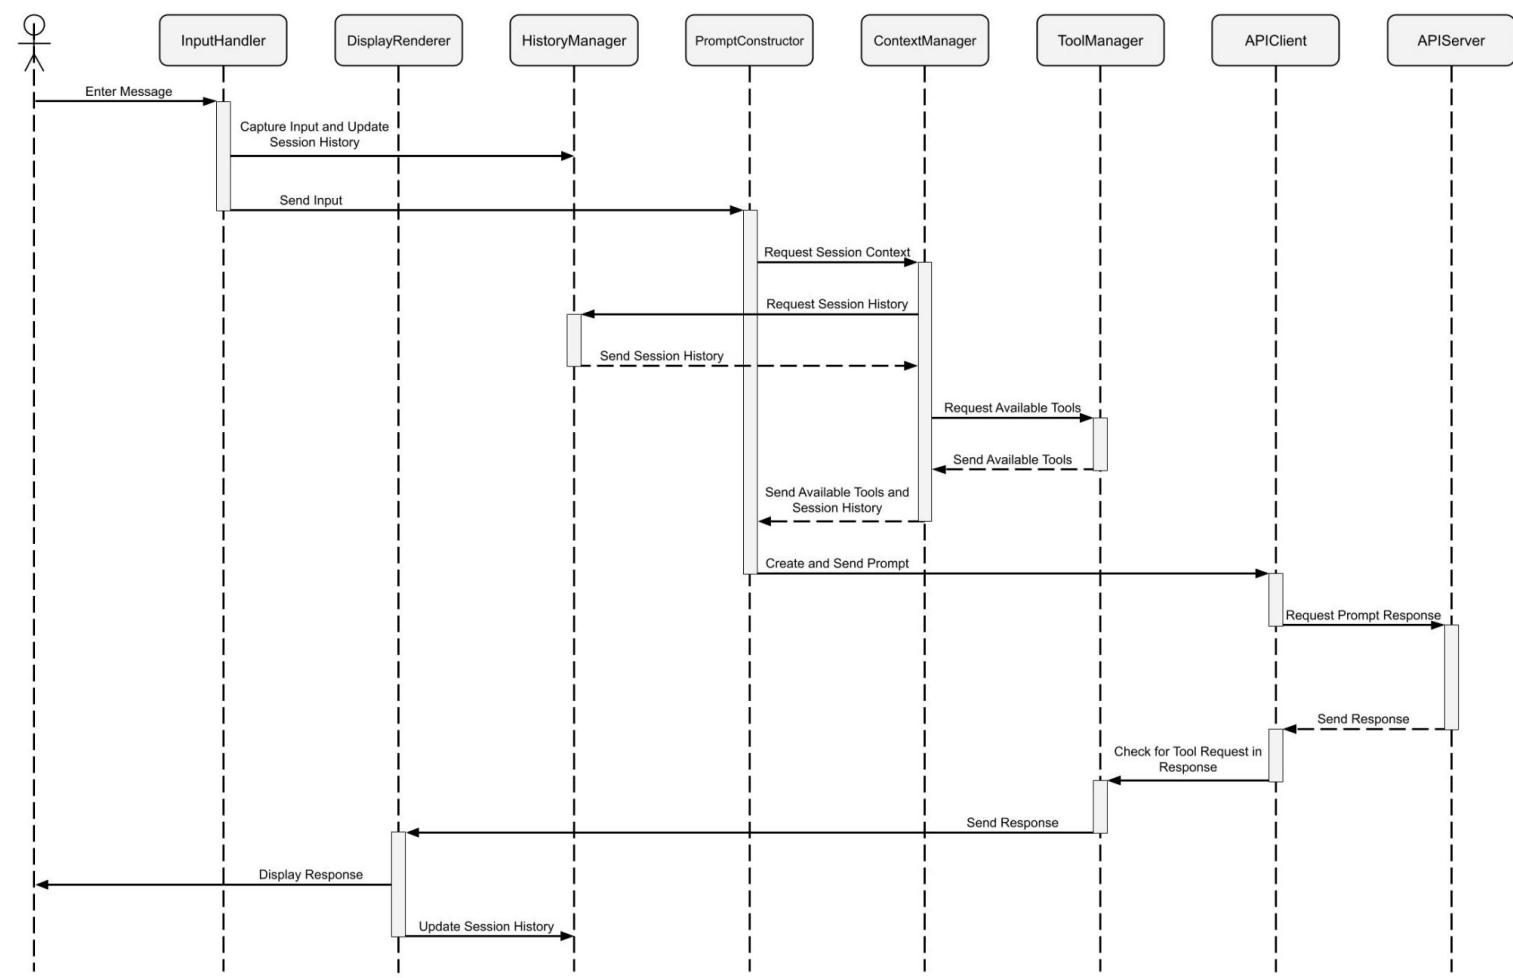
Sub-component



Sequence Diagram

User Prompt Without
Tool Execution:





Legend



Actor



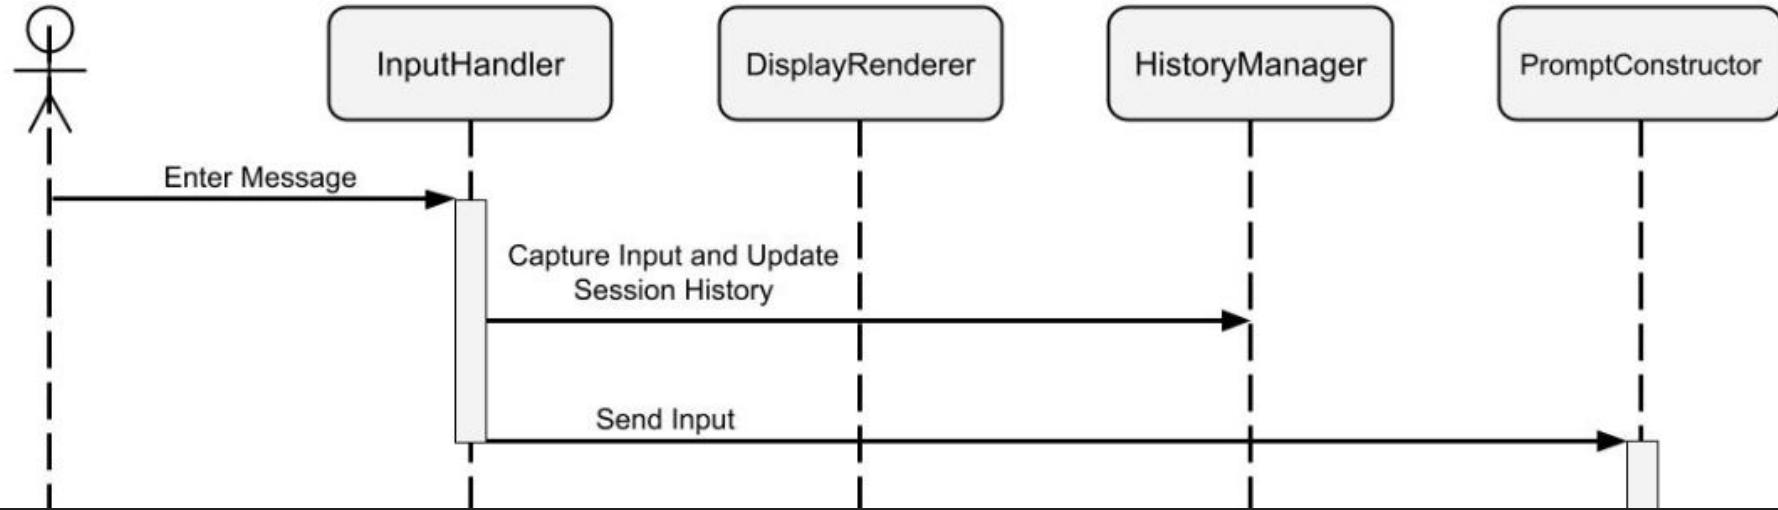
Component

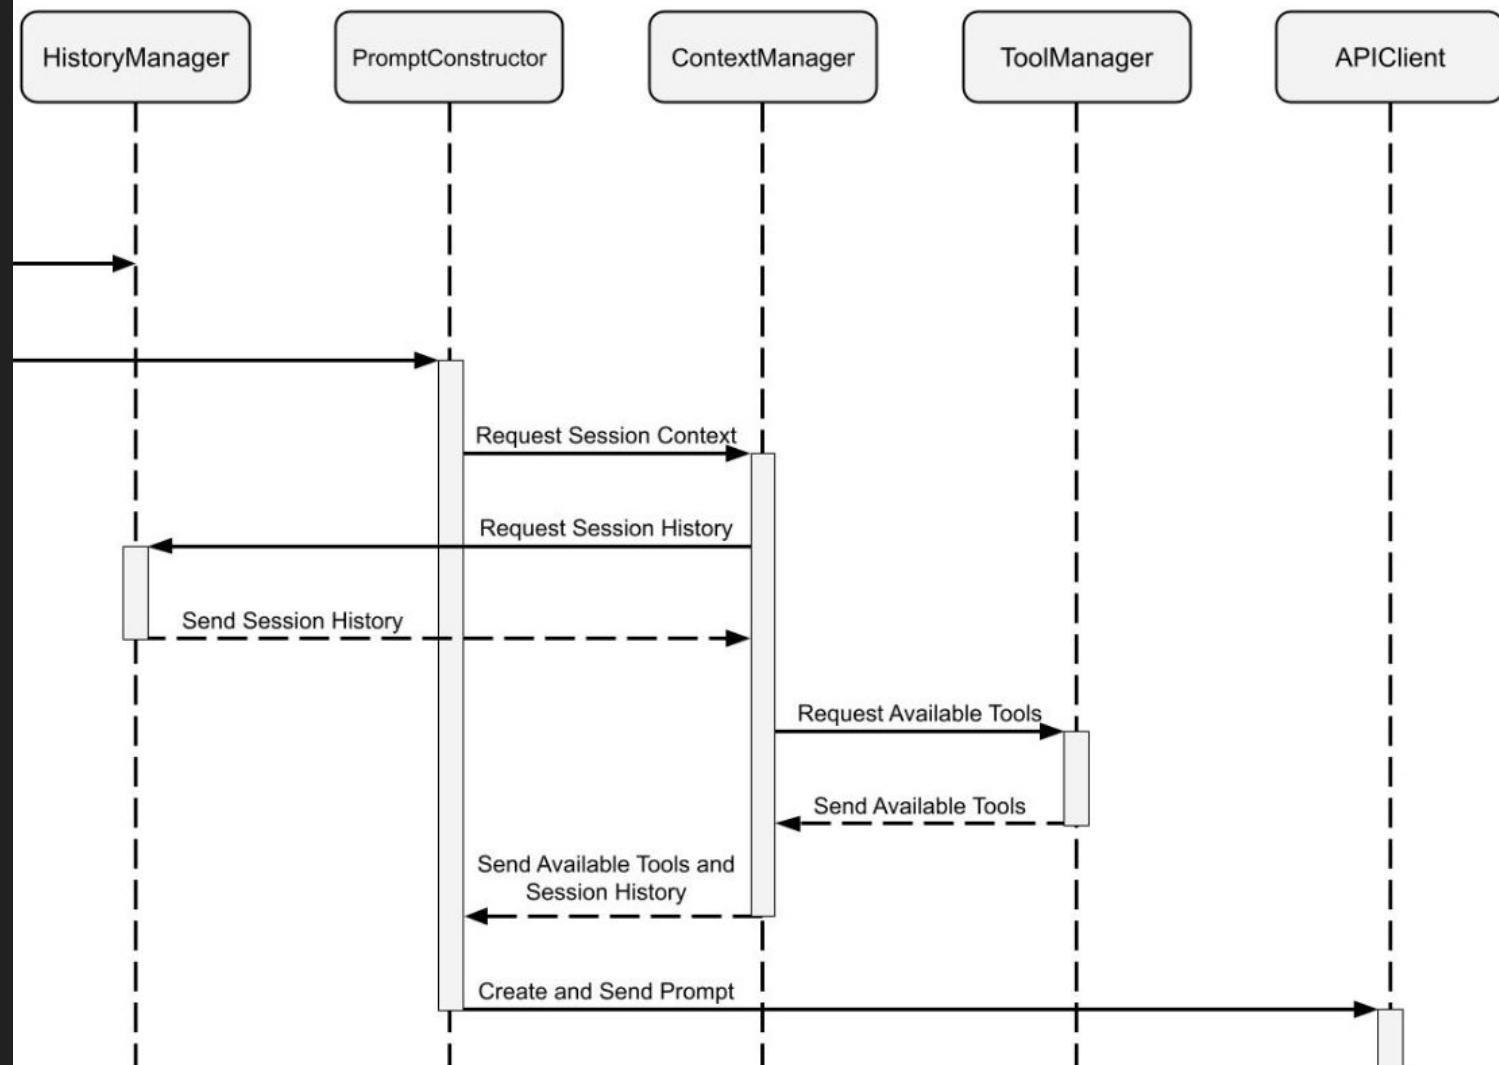


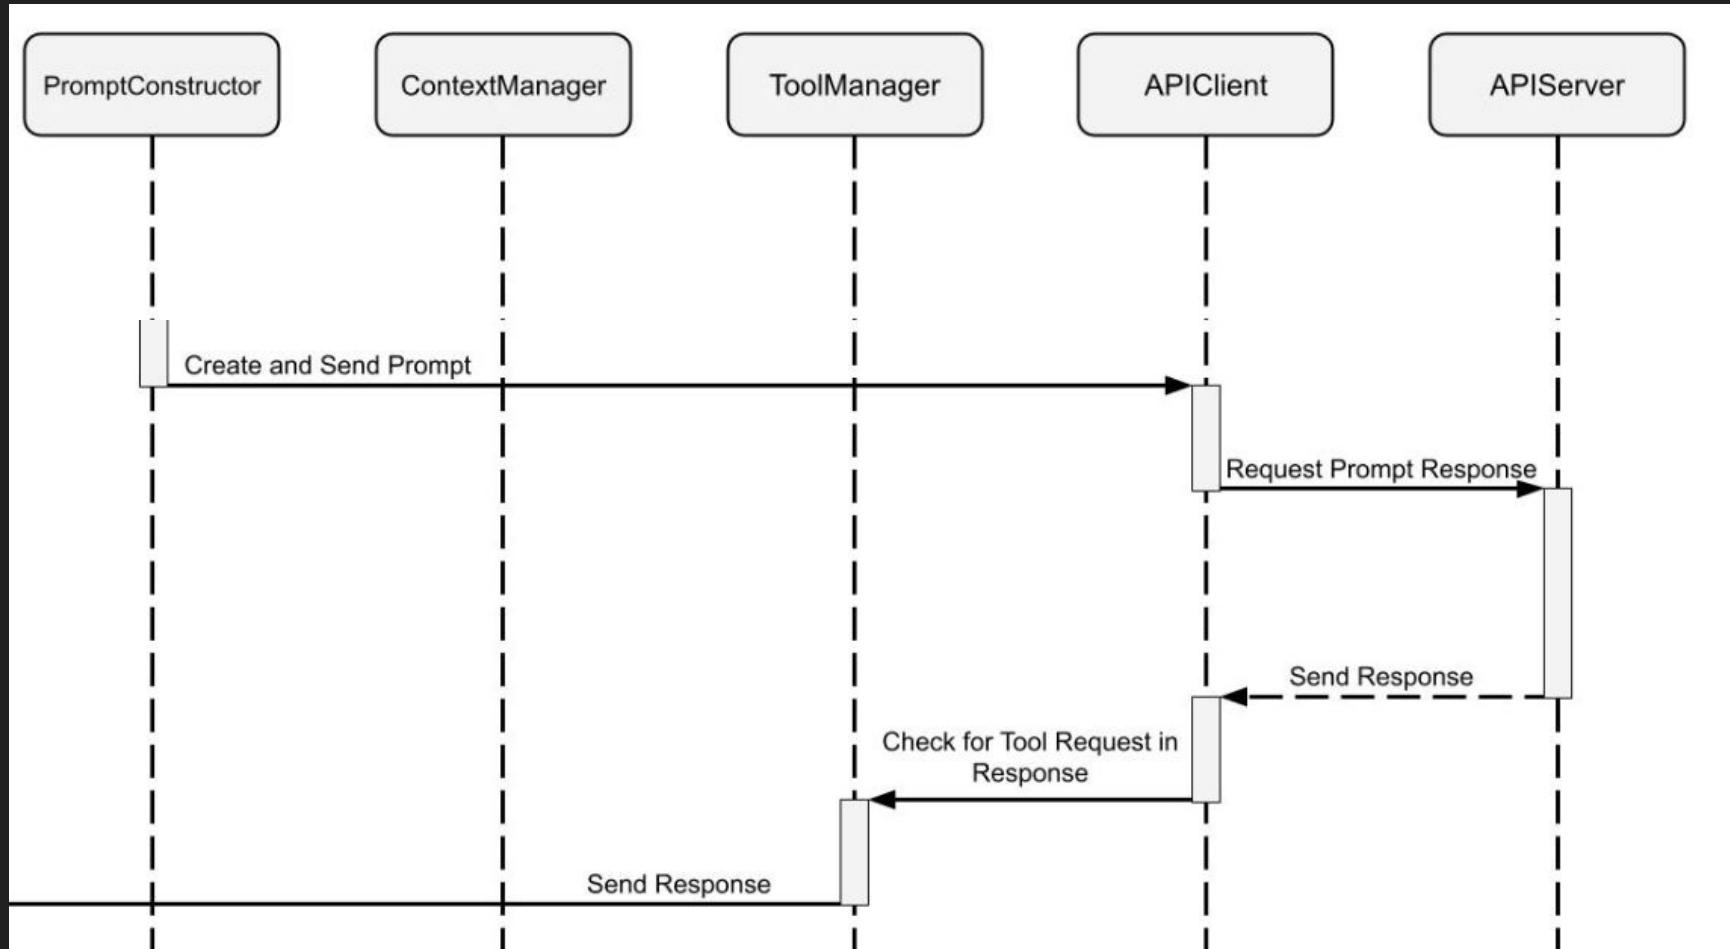
Activation Bar

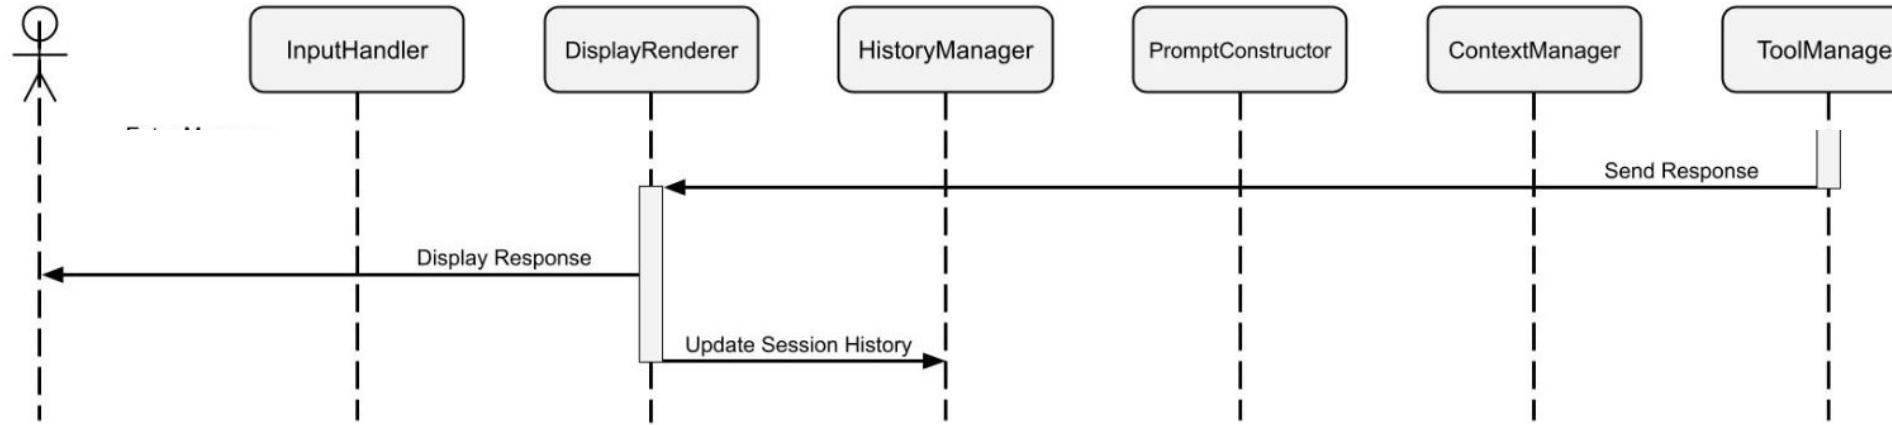
Message

Response











Interactions Between Subsystems

- **User Input:** InputManager → HistoryManager + PromptConstructor
- **State Recovery:** ContextManager → HistoryManager
- **Context Constructing:** PromptConstructor → ContextManager
- **Prompt Sending:** APIClient → PromptConstructor
- **Response Interpreting:** ToolManager → APIClient
- **Human-In-The-Loop:** ToolManager → User/DisplayRenderer
- **Remote Comms:** API Server → APIClient
- **Tool Execution:** Tools → ToolManager



Derivation Process

1. Gemini CLI Documentation
 - Layered and Client / Server
2. Gemini CLI Demo Video & Use Case Diagrams
 - User Entering a Prompt
3. Box-and-Arrow Diagram
 - Dependencies between Components



Alternative Architectures

Pipe and Filter

- System viewed as series of independent transformations
- Input = data stream
- Goes through core then to Gemini API
- Improves concurrency
- Less interactive, no back and forth discourse
- Overall less efficient

Interpreter

- Core = execution engine
- Gemini API = instruction set
- CLI takes human texts, makes it executable actions
- Follows interpreter process
- CLI abstracts LLM logic
- Adds much latency
- Instruction set is non deterministic



Lessons Learned

Necessity of Strict Interfaces

- Maintaining boundaries between CLI, Core, API were paramount
- Prevents blurring between layers

Asynchronous Requirements Add Complexity

- Event driven model important, but makes control flow more complexed
- ReAct loop pauses and resumes many times



Lessons Learned (cont'd)

Safety Must be a Priority

- “Human-in-the-loop” fundamental requirement
- Since LLMs are non deterministic, human review is required
- Security needs to be implemented deep within

Visualization Driven Discovery

- Took creating the sequence diagram to understand components
- Visualizing data flow revealed missing dependencies
- Allowed us to add ContextManager, HistoryManager



Limitations of Findings

- All similar to the Lessons Learnt slide
- Limited understanding of inner workings of API
- Potentially other components missing from sequence diagram
- Impossible to map the AI behaviour due to non-deterministic nature



AI Collaboration

LLMs used: GPT-4, Claude 4.5, Gemini 3 Pro

- Using Gemini for Gemini CLI was important (up to date)

Used For:

- Rubric Compliance
- Architectural Analysis
- Technical Articulation



AI Collaboration

- Used 'driver-navigator' approach - designated person to manage the interactions. Helped with consistent workflow and consistent prompts
- Prompting Strategy: word limits, specifying formatting. Injecting more information
- Verification Strategy: checked claims with CLI documentation. Made sure all information matched
- Attributed to ~25% of the final deliverable
- Streamlined workflow, improving pace
- Sometimes would hallucinate features